

Rubrik

Integrating with ELK

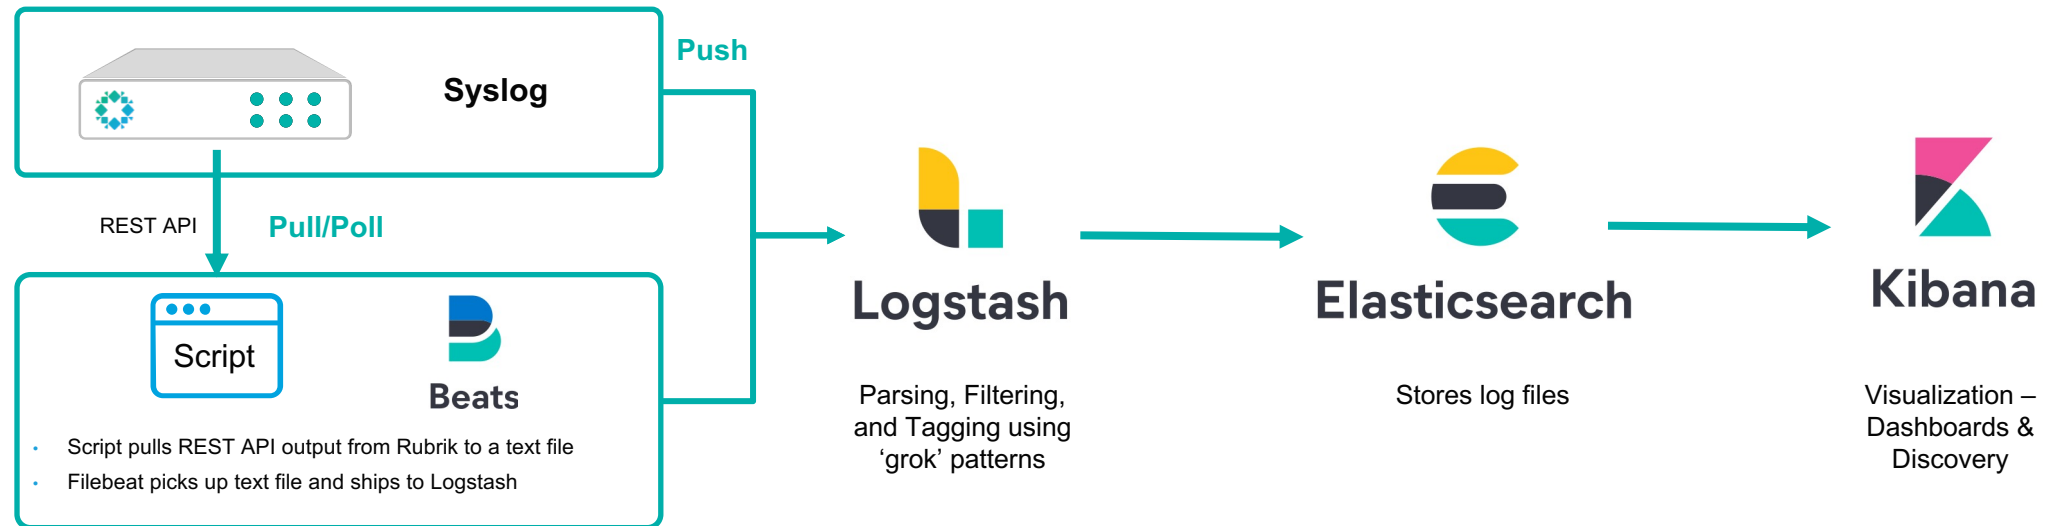
Steven Tong
February 2021



Monitoring Rubrik

- Can Rubrik integrate with <ELK, Splunk, LogRhythm, SolarWinds, Datadog, log management/SIEM application of your choice>?
 - Yes, Rubrik can send (push) Syslog events or SNMP traps to the log/SIEM app or you can pull/poll for info using our REST API
 - In v5.3 we can also send Syslog events as SNMP traps – requires installing our three MIBs (download MIBs from the SNMP config screen in the CDM UI)
- What type of events does Rubrik send?
 - Things like job activity, audit events, SLA changes, hardware failures
- What if I wanted to build dashboards for things like capacity and SLA compliance?
 - That is when you would use the REST API to pull/poll for information to bring in any additional info you want

Rubrik for ELK



<https://github.com/stevenctong/rubrik/tree/main/elk>

ELK for Syslog

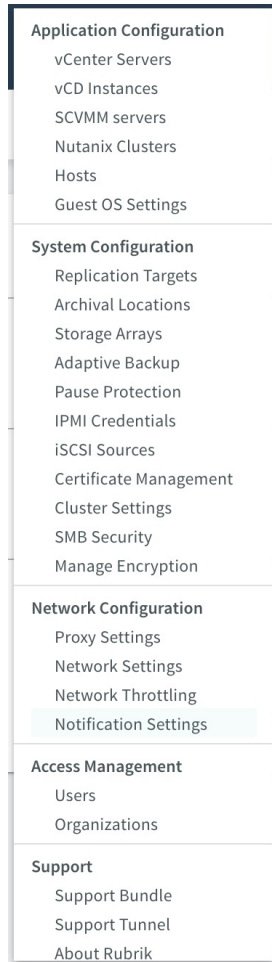
1. **Rubrik** – Configure Syslog to send events to Logstash
2. **Logstash** – Configure .conf files to parse Syslog messages to send to an Elasticsearch index
 - /etc/logstash/conf.d/
3. **Kibana** – Create an ‘index pattern’ in order to use the Elasticsearch index
4. **Kibana** – Discover Syslog data and create dashboards

ELK for Monitoring Status / Metrics

1. **Script Host** – Configure and schedule the script to poll REST API information at an interval and write the output to a log file in JSON format
 - `get_rubrik_stats.sh + rubrik_cluster.conf` – Bash script
 - `get_rubrik_stats.ps1` – Powershell script
2. **Script Host** – Configure Filebeat to monitor the log file folder and send any updates to Logstash for further processing
 - `/var/log/rubrikelk/*.log`
3. **Logstash** – Configure `.conf` file to format fields to send to an Elasticsearch index
 - `/etc/logstash/conf.d/`
4. **Kibana** – Create an 'index pattern' in order to use the Elasticsearch index
5. **Kibana** – Discover the data and create dashboards

Note: We're using Filebeat (logs) and not Metricbeat (metrics) so the data in Elasticsearch is log-based and not metric-based but we can still create the dashboards we want

Rubrik – Adding a Syslog Rule



Email Settings SNMP Syslog Notifications							Add Syslog Export Rule	
IP or Hostname	Protocol	Port	Facility	Severity	TLS			
gaia-splunk01.rubrikdemo.com	UDP	514	RubrikEvent	Informational	Disabled	...		
sx1-stevtong-l1.rubrikdemo.com	UDP	5000	RubrikEvent	Informational	Disabled	...		

Add Syslog Export Rule

IP or Hostname
<syslog host>

Protocol
☐ TCP ☒ UDP

Port
5000

Facility
RubrikEvent

Severity
Informational

☐ Enable TLS

Cancel [Add](#)

- **Gear → Notification Settings**
- **Syslog tab**
- **Add Syslog Export Rule button**
- Fill in IP or hostname of your Logstash host, protocol (TCP or UDP), and port #
- Facility “RubrikEvent” and Severity “Informational” should be fine or change as you see fit

- Logstash configuration files are used to configure Inputs → Filters → Output
 - Default directory –
/etc/logstash/conf.d
- **Inputs** – Monitors the TCP & UDP ports for incoming Syslog messages
 - Syslog messages are received over non-privileged ports and then parsed
- **Filter** – Uses grok to parse the Syslog messages to fields
- **Output** – Fields are sent to Elasticsearch & any grok failures are also logged
 - Sends to a specific 'index' in Elasticsearch

Rubrik Syslog Examples

```
<134>1 2021-01-30T19:13:35.672Z RVM16CS014306 Rubrik-SprayServer 31932 Rubrik [eventDetail@49929 clusterName=\"sand1-rbk01\" errorCode=\"-\" errorMessage=\"-\" errorReason=\"-\" errorRemedy=\"-\" eventId=\"884cc6df-293b-4d30-9b17-b62d1a0b3d56\" eventName=\"Audit.CreatedManagedVolumeSnapshotAudit\" eventSeriesId=\"786a194d-1df6-44f6-90ea-0ec28ea59469\" eventSeverity=\"Informational\" eventType=\"Audit\" locationName=\"-\" nodeId=\"RVM16CS014306\" nodeIpAddress=\"172.xx.yy.zz\" objectId=\"15bda393-bd81-4f1c-b731-5ae54f9500e2\" objectName=\"-\" objectType=\"UserActionAudit\" status=\"Success\"] saphana_sh1-hxexsa-2_HXE created a snapshot 8ca08bdc-ea5c-48c1-8336-02385cf192fa for managed volume 'sap_hana_sh1-hxexsa-2_HXE_SYSTEMDB_log'\n"
```

```
<134>1 2021-01-30T17:34:13.229Z RVM16CS014476 Rubrik-JobFetcherLoop 8833 Rubrik [eventDetail@49929 clusterName=\"sand1-rbk01\" errorCode=\"-\" errorMessage=\"-\" errorReason=\"-\" errorRemedy=\"-\" eventId=\"f4f549fb-4a3b-4eae-837e-3e97d87cd708\" eventName=\"Snapshot.LogBackupSucceeded\" eventSeriesId=\"2a517e35-4d40-4a57-a698-bc3174a602bd\" eventSeverity=\"Informational\" eventType=\"Backup\" locationName=\"Oracle Host sx1-shawmcel-l3\" nodeId=\"RVM16CS014476\" nodeIpAddress=\"172.xx.yy.zz\" objectId=\"d805203a-3eb4-4289-bf7b-10505b41e794\" objectName=\"smcdb01\" objectType=\"OracleDb\" status=\"Success\"] [mdc@49929 JobContextQueryStatsSnapshotId=\"1731619182330813365\" instanceId=\"-1\" jobId=\"\" jobType=\"\" ndc=\"CREATE_ORACLE_LOG_SNAPSHOT_d805203a-3eb4-4289-bf7b-10505b41e794:::906\" parentSpanId=\"0\" profile=\"false\" spanId=\"313acb5790465439\" taskId=\"\" tracerId=\"b97fdd05fcaae4d6313acb5790465439\"] Completed log backup of Oracle Database 'smcdb01'\n"
```

- Rubrik Syslog messages follow a standard format
- One of the formats has an extra array so two 'grok' rules will need to be created to parse the messages

Parsing With Grok

```
<131>1 2021-02-16T07:42:54.457Z RVM16CS014476 Rubrik-JobFetcherLoop 8833 Rubrik [eventDetail@49929 clusterName="sand1-rbk01" errorCode="RBK20700003" errorMessage="Unable to resolve host name 'sx1-test-w1.rubrikdemo.com'." errorReason="Failed to look up the IP address for host 'sx1-test-w1.rubrikdemo.com'." errorRemedy="Make sure correct host name is used and DNS server is configured correctly." eventId="a54c7b50-70b2-48e2-9de6-930a3f104c10" eventName="Snapshot.BackupFailed" eventSeriesId="5a39f449-7b31-42fc-822c-6dd8653cace0" eventSeverity="Warning" eventType="Backup" locationName="-" nodeId="RVM16CS014476" nodeIpAddress="172.xx.yy.zz" objectId="f6d37b0f-32fa-49af-b543-4d96f3e0e0e9" objectName="sx1-test-w1.rubrikdemo.com volumes" objectType="VolumeGroup" status="Failure"] [mdc@49929 JobContextQueryStatsSnapshotId="3510858650583934103" instanceId="-1" jobId="" jobType="" ndc="CREATE_VOLUME_GROUP_SNAPSHOT_f6d37b0f-32fa-49af-b543-4d96f3e0e0e9:::8068" parentSpanId="0" profile="false" spanId="966367648eb0bf18" taskId="" tracerId="e95fc547e0d25c2966367648eb0bf18"] Failed backup of Volume Group 'sx-test-w1.rubrikdemo.com volumes'. Reason: RBK20700003 - Unable to resolve host name 'sx1-test-w1.rubrikdemo.com'.
```

Grok statements that parse the above example

- %{TIMESTAMP_ISO8601:timestamp} - 2021-02-16T07:42:54.457Z
- %{NOTSPACE:rubrikNodeSN} - RVM16CS014476
- %{NOTSPACE:rubrikComponent} - Rubrik-JobFetcherLoop
- %{NOTSPACE:rubrikNum} - 8833
- %{NOTSPACE} - Rubrik
- %{GREEDYDATA} clusterName=\ - [eventDetail@49929 clusterName=
- %{NOTSPACE:rubrikClusterName} - sand1-rbk01
- \ " errorCode=\ - " errorCode=
- %{GREEDYDATA:rubrikErrorCode} - RBK20700003
- \ " errorMessage=\ - " errorMessage=
- %{GREEDYDATA:rubrikErrorMessage} - Unable to resolve host name 'sx1-test-w1.rubrikdemo.com'.
- \ " errorReason=\ - " errorReason=
- %{GREEDYDATA:rubrikErrorReason} - Failed to look up the IP address for host 'sx1-test-w1.rubrikdemo.com'.

...

Some 'grok' resources

- <https://grokdebug.herokuapp.com/>
- https://streamsets.com/documentation/datacollector/latest/help/datacollector/UserGuide/Apx-GrokPatterns/GrokPatterns_title.html
- <https://www.kartar.net/2014/09/when-logstash-and-syslog-go-wrong/>
- **Green** – The field + message that is stored in the matched pattern
- **Red** – Matched pattern that is not stored in a field (thrown away0)

Scripts & Filebeat

- A script can be scheduled to run periodically to gather status and metrics from Rubrik
 - get_rubrik_stats.sh (bash)
 - get_rubrik_stats.ps1 (Powershell)
- The data is captured in JSON format to a log file
- Filebeat is configured to look for updates to those log files in a certain directory and send it to Logstash for further processing
 - /var/log/rubrikelk/*.log

filebeat.yml

```
##### Filebeat Configuration Example #####

# ===== Filebeat inputs =====

filebeat.inputs:

- type: log

# Change to true to enable this input configuration.
enabled: true

# Paths that should be crawled and fetched. Glob based paths.
paths:
  #- /var/log/*.log
  - /var/log/rubrikelk/*.log

# In the event the Filebeat cannot send to output, close open files after timeout
# Data loss is a potential side effect
close_timeout: 1h

# ----- Logstash Output -----
output.logstash:
  # The Logstash hosts
  hosts: ["localhost:5044"]
```

Logstash (Filebeat)

- **Inputs** – Uses 'beats' filter
 - Beats messages are received over non-privileged port and then parsed
- **Filter** – The data being sent should be in JSON format
 - Re-name fields as needed
 - Whitelist the fields that will be sent to Elasticsearch
- **Output** – Elasticsearch

20-logstash-rubrik-filebeat.conf

```
input {
  beats {
    port => 5044
    type => "beats"
  }
}

filter {
  if [type] == "beats" {

    json {
      source => "message"
    }
    mutate {
      rename => [ "lastUpdateTime", "timestamp" ]
    }

    date {
      match => [ "scriptRunTime" , "ISO8601" ]
      target => "@timestamp"
    }
    prune {
      whitelist_names => [ "timestamp", "type", "index", "rubrikClusterName", "rubrikSpaceTotal",
        "rubrikSpaceUsed", "rubrikSpaceAvailable", "rubrikSpaceSnapshot", "rubrikSpaceLiveMount",
        "rubrikSpacePendingSnapshot", "rubrikSpaceCDP", "rubrikSpaceMisc", "rubrikUsedPct",
        "rubrikTotalProtected", "rubrikInCompliance", "rubrikOutCompliance",
        "rubrikPctInCompliance", "rubrikPctOutCompliance", "rubrikComplianceTime",
        "rubrikNodesGood", "rubrikNodesBad", "rubrikNodesTotal" ]
    }
  }
}

output {
  if [type] == "beats" {
    elasticsearch {
      hosts => ["localhost:9200"]
      index => "logstash-rubrik-filebeat"
    }

    stdout { codec => rubydebug }
  }
}
```

Kibana – Indexes

Elastic Search Elastic

Stack Management / Index Management

Ingest ②
Ingest Node Pipelines

Data ②
Index Management
Index Lifecycle Policies
Snapshot and Restore
Rollup Jobs
Transforms
Remote Clusters

Alerts and Insights ②
Alerts and Actions
Reporting

Kibana ②

Index Management

[Index Management docs](#)

[Indices](#) Data Streams Index Templates Component Templates

Update your Elasticsearch indices individually or in bulk. [Learn more.](#) ☐ Include rollup indices ☐ Include hidden indices

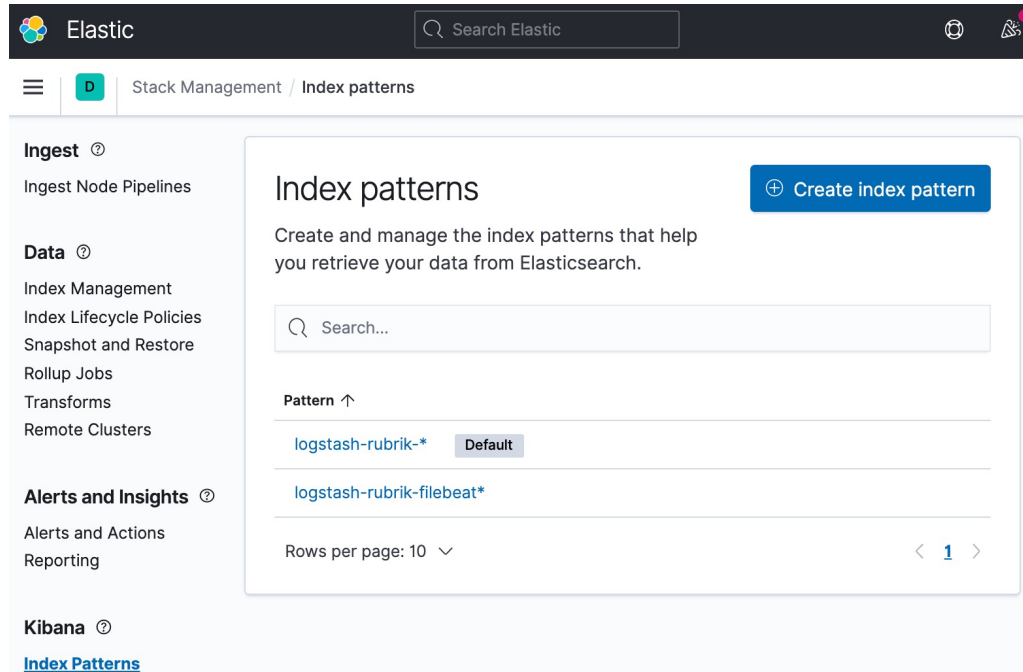
Lifecycle status Lifecycle phase [Reload indices](#)

<input type="checkbox"/> Name	Health	Status	Primaries	Replicas	Docs count	Storage size	Data stream
<input type="checkbox"/> logstash-rubrik-syslog	● yellow	open	1	1	2068420	1gb	
<input type="checkbox"/> logstash-rubrik-filebeat	● yellow	open	1	1	4360	1.4mb	

Rows per page: 10 < 1 >

- Data sent to Elasticsearch indexes should start appearing under Stack Management → Index Management
- Log lifecycle and management will need to be configured in order to manage the size of the indexes

Kibana – Index Patterns



The screenshot shows the Kibana interface. The top navigation bar includes the Elastic logo, a search bar, and user profile icons. The left sidebar contains navigation links for Ingest, Data, Alerts and Insights, and Kibana. The main content area is titled 'Index patterns' and includes a 'Create index pattern' button. Below this is a search bar and a table of existing patterns. The table has two rows: 'logstash-rubrik-*' (marked as 'Default') and 'logstash-rubrik-filebeat*'. At the bottom of the table, it says 'Rows per page: 10' and shows a pagination control for page 1.

Elastic

Stack Management / Index patterns

Ingest

Ingest Node Pipelines

Data

Index Management
Index Lifecycle Policies
Snapshot and Restore
Rollup Jobs
Transforms
Remote Clusters

Alerts and Insights

Alerts and Actions
Reporting

Kibana

[Index Patterns](#)

Index patterns

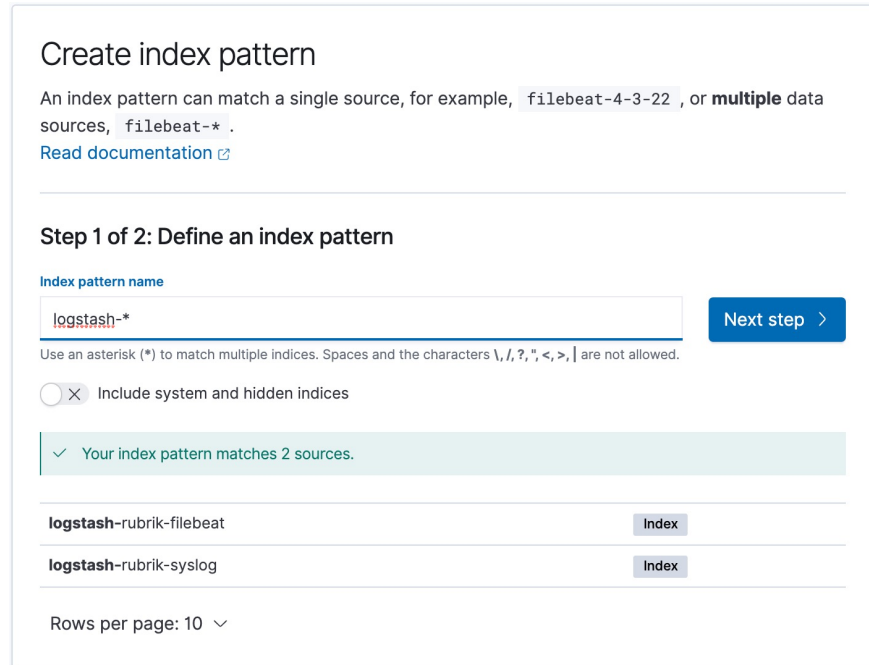
Create and manage the index patterns that help you retrieve your data from Elasticsearch.

Search...

Pattern ↑

logstash-rubrik-*	Default
logstash-rubrik-filebeat*	

Rows per page: 10



The screenshot shows the 'Create index pattern' wizard. It starts with an explanation of index patterns and a link to the documentation. The current step is 'Step 1 of 2: Define an index pattern'. It features a text input for the 'Index pattern name' with the value 'logstash-*'. Below the input is a note about using asterisks and a checkbox for 'Include system and hidden indices'. A green success message states 'Your index pattern matches 2 sources.' Below this is a table showing the matching indices: 'logstash-rubrik-filebeat' and 'logstash-rubrik-syslog', both with an 'Index' label. At the bottom, it says 'Rows per page: 10'.

Create index pattern

An index pattern can match a single source, for example, `filebeat-4-3-22`, or **multiple** data sources, `filebeat-*`.

[Read documentation](#)

Step 1 of 2: Define an index pattern

Index pattern name

logstash-*

Use an asterisk (*) to match multiple indices. Spaces and the characters `\, /, ?, ", <, >, |` are not allowed.

☐ Include system and hidden indices

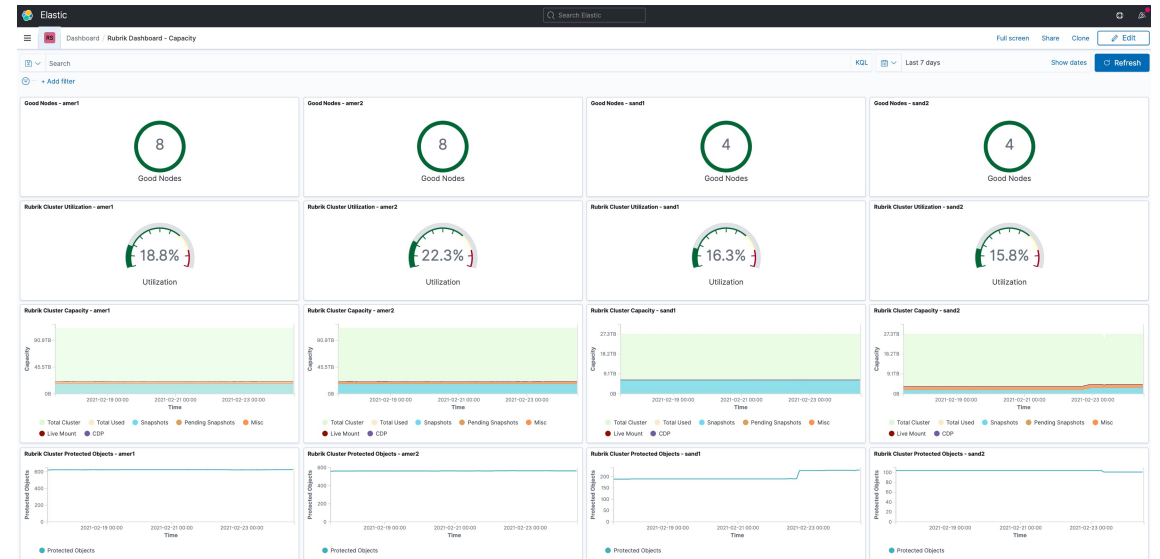
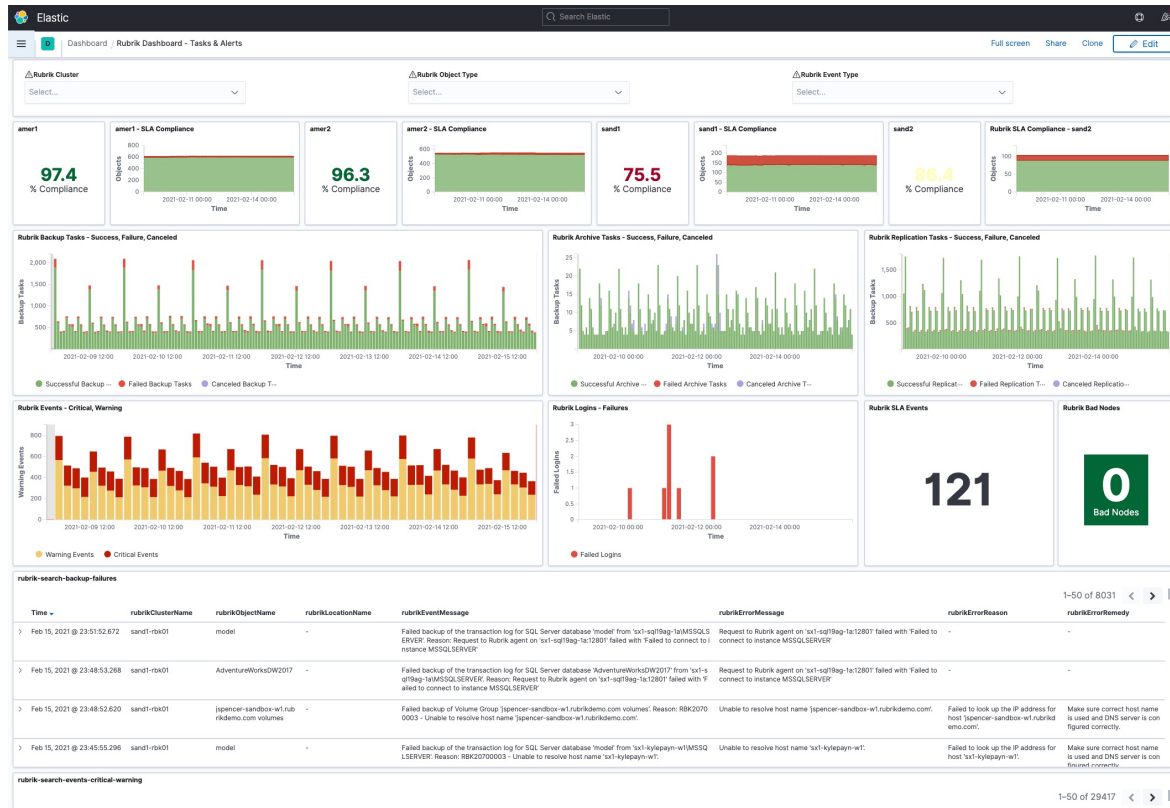
✓ Your index pattern matches 2 sources.

logstash-rubrik-filebeat	Index
logstash-rubrik-syslog	Index

Rows per page: 10

- Create an Index Pattern under Stack Management → Index Patterns
- One or more Index Patterns must be created in order to reference data stored in the Elasticsearch indexes
- When you create a pattern name it will let you know all indexes that match that pattern

Kibana – Dashboards



- Two example dashboards are created in Kibana
- For all dashboards you can change the time range and provide additional search filters to update the results on the dashboard

Dashboard - Tasks & Alerts

All tiles were created from Syslog events except for SLA Compliance which uses metrics polled from the REST API

Charts are preferable when possible for trend analysis

- **Drop-down Filter** – Filter the data shown by cluster, object type, and event type
- **Number + Chart – SLA Compliance** – For each cluster display as a % and chart
- **Chart – Backup Tasks** – Display successful, failed, and cancelled backup tasks
- **Chart – Archive Tasks** – Display successful, failed, and cancelled archive tasks
- **Chart – Replication Tasks** – Display successful, failed, and cancelled replication tasks
- **Chart – Events (Critical, Warning)** – Display events that are Warning or Critical level
- **Chart – Login Failures** – Display number of login failures
- **Number – SLA Failures** – Display a # of SLA type events (ie changes) over the current time range
- **Number – Bad Nodes** – Display the number of bad nodes currently seen against a colored background
- **Event Log – Backup Failures** – List failed backup events
- **Event Log – Critical, Warning** – List events that are Warning or Critical level
- **Event Log – SLA** – List events that modified SLA domains
- **Event Log – Failed Logins** – List events of failed logins

Dashboard - Capacity

All tiles were created using REST API metrics

Aggregate metrics are hard to create and do not make sense so separate tiles are created for each cluster

- **Number + Donut – Good Nodes** – Display the number of good nodes along with a colored, filled donut of the number of good nodes
- **Number + Gauge – Cluster Utilization** – Show current cluster utilization within a colored gauge that goes up to 100%
- **Chart – Cluster Capacity** – Display the cluster storage utilization by TB over time with capacity breakouts
 - It can be hard to hover over this chart – “Inspect” the chart to see the actual data values
- **Chart – Protected Objects** – Display number of protected objects over time

Don't Backup.
Go Forward.

